

u-s-home-assignment-final

December 13, 2023

Data Science Project: Analyzing Key Factors Influencing US Home Prices

Introduction

The data for this analysis was collected from various publicly available sources, focusing on key factors believed to influence home prices in the United States. Using techniques such as VLOOKUP and HLOOKUP in Excel, the collected data was merged to create a comprehensive dataset. This dataset is now imported into this Colab notebook for further analysis and exploration.

Objective

The primary objective of this data science project is to build a predictive model that explains how different factors have impacted home prices in the United States over the last two decades (20 years). The analysis will revolve around understanding the relationship between various economic, demographic, and housing market indicators and the fluctuations in home prices.

```
[2]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline
from sklearn.metrics import r2_score, mean_squared_error, mean_absolute_error
```

```
[3]: df = pd.read_csv("/content/Assignment_1.csv")
```

```
[4]: df.head(1)
```

```
[4]:      date  new_constructed_units  unemployment_rate  interest_rates \
0  01-01-1947                    NaN                    NaN            NaN

      income  home_price_index  per_capita_gdp  urban_population
0      NaN                NaN          15248                NaN
```

```
[5]: df.tail(1)
```

```
[5]:          date  new_constructed_units  unemployment_rate  interest_rates  \
921  01-10-2023                1410.0                3.9                5.33

      income  home_price_index  per_capita_gdp  urban_population
921  16848.7            311.175            67083            83.084
```

```
[6]: df.describe().T
```

```
[6]:          count          mean          std          min  \
new_constructed_units  670.0    1396.762687    358.271084    520.000
unemployment_rate      910.0         5.709670         1.708239         2.500
interest_rates         832.0         4.599916         3.596581         0.050
income                 778.0    8255.129820    4299.805894    2318.400
home_price_index       442.0     142.293441         61.649943         63.965
per_capita_gdp         922.0    37430.843818    15402.352634    15032.000
urban_population       766.0         76.764057         3.788336         69.996

          25%          50%          75%          max
new_constructed_units  1190.00000    1396.000    1641.7500    2299.000
unemployment_rate      4.40000         5.500         6.7000     14.700
interest_rates         1.79000         4.160         6.2425     19.100
income                4539.62500    7282.850    12069.1750    20422.600
home_price_index       82.25075      141.275      179.0250     311.175
per_capita_gdp        24176.75000    35582.000    52835.0000    67083.000
urban_population       73.65550         75.701         80.2690     83.084
```

```
[7]: df.nunique()
```

```
[7]: date                922
new_constructed_units    536
unemployment_rate        83
interest_rates          503
income                  777
home_price_index        440
per_capita_gdp          307
urban_population        63
dtype: int64
```

```
[8]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 922 entries, 0 to 921
Data columns (total 8 columns):
#   Column                Non-Null Count  Dtype
---  -
0   date                  922 non-null   object
1   new_constructed_units  670 non-null   float64
```

```

2    unemployment_rate      910 non-null    float64
3    interest_rates         832 non-null    float64
4    income                 778 non-null    float64
5    home_price_index       442 non-null    float64
6    per_capita_gdp         922 non-null    int64
7    urban_population       766 non-null    float64
dtypes: float64(6), int64(1), object(1)
memory usage: 57.8+ KB

```

```
[9]: df['date'] = pd.to_datetime(df['date'])
```

```
[10]: df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 922 entries, 0 to 921
Data columns (total 8 columns):
#   Column                Non-Null Count  Dtype
---  -
0   date                  922 non-null   datetime64[ns]
1   new_constructed_units  670 non-null   float64
2   unemployment_rate     910 non-null   float64
3   interest_rates        832 non-null   float64
4   income                778 non-null   float64
5   home_price_index      442 non-null   float64
6   per_capita_gdp        922 non-null   int64
7   urban_population      766 non-null   float64
dtypes: datetime64[ns](1), float64(6), int64(1)
memory usage: 57.8 KB

```

```
[11]: df.head(1)
```

```

[11]:      date  new_constructed_units  unemployment_rate  interest_rates \
0  1947-01-01                    NaN                    NaN                    NaN

      income  home_price_index  per_capita_gdp  urban_population
0      NaN                NaN            15248                NaN

```

```
[12]: df_20 = df[(df["date"] >= '01-01-2003') & (df["date"] <= '01-10-2023')]
```

```
[13]: df_20.head(1)
```

```

[13]:      date  new_constructed_units  unemployment_rate  interest_rates \
672  2003-01-01            1654.0                5.8                1.24

      income  home_price_index  per_capita_gdp  urban_population
672  10710.4            128.461            50462            79.583

```

```
[14]: df_20.tail(1)
```

```
[14]:      date  new_constructed_units  unemployment_rate  interest_rates  \
921  2023-01-10                1410.0                3.9                5.33

      income  home_price_index  per_capita_gdp  urban_population
921  16848.7             311.175           67083             83.084
```

```
[15]: df_20.isnull().sum()
```

```
[15]: date                0
new_constructed_units    0
unemployment_rate        0
interest_rates           0
income                   0
home_price_index         0
per_capita_gdp           0
urban_population         0
dtype: int64
```

There are no Null Values on our dataset.

```
[16]: df_20.duplicated().sum()
```

```
[16]: 0
```

There are no Duplicate columns present in our dataset.

```
[17]: df_20.nunique()
```

```
[17]: date                250
new_constructed_units    228
unemployment_rate        64
interest_rates           118
income                   250
home_price_index         249
per_capita_gdp           83
urban_population         20
dtype: int64
```

Above info shows that there are no categorical variables present in the dataset.

```
[18]: df_20['year'] = df_20['date'].dt.year
```

```
<ipython-input-18-b94521149b84>:1: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: <https://pandas.pydata.org/pandas->

docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
df_20['year'] = df_20['date'].dt.year
```

Creating a new variable “Year” to make Year wise Analysis for the variables.

```
[19]: df_20.head()
```

```
[19]:
```

	date	new_constructed_units	unemployment_rate	interest_rates	\
672	2003-01-01	1654.0	5.8	1.24	
673	2003-01-02	1688.0	5.9	1.26	
674	2003-01-03	1638.0	5.9	1.25	
675	2003-01-04	1662.0	6.0	1.26	
676	2003-01-05	1733.0	6.1	1.26	

	income	home_price_index	per_capita_gdp	urban_population	year
672	10710.4	128.461	50462	79.583	2003
673	10674.0	129.355	50462	79.583	2003
674	10696.5	130.148	50462	79.583	2003
675	10752.7	130.884	50796	79.583	2003
676	10832.0	131.735	50796	79.583	2003

```
[20]: pd.set_option('display.max_columns',None)
df_20.
↳groupby('year')['new_constructed_units','unemployment_rate','interest_rates','income','home
↳agg(['mean', 'max', 'min']).T
```

<ipython-input-20-7ada3471925c>:2: FutureWarning: Indexing with multiple keys (implicitly converted to a tuple of keys) will be deprecated, use a list instead.

```
df_20.groupby('year')['new_constructed_units','unemployment_rate','interest_ra
tes','income','home_price_index','per_capita_gdp','urban_population'].agg(['mean
','max','min']).T
```

```
[20]:
```

year		2003	2004	2005	\
new_constructed_units	mean	1676.750000	1834.583333	1929.333333	
	max	1733.000000	1938.000000	2103.000000	
	min	1570.000000	1709.000000	1787.000000	
unemployment_rate	mean	5.991667	5.541667	5.083333	
	max	6.300000	5.800000	5.400000	
	min	5.700000	5.400000	4.900000	
interest_rates	mean	1.127500	1.349167	3.213333	
	max	1.260000	2.160000	4.160000	
	min	0.980000	1.000000	2.280000	
income	mean	10884.391667	11233.175000	11364.858333	
	max	11066.700000	11659.000000	11535.200000	
	min	10674.000000	11051.200000	11226.500000	
home_price_index	mean	133.731333	150.440250	171.737000	
	max	140.179000	159.330000	180.910000	

	min	128.461000	141.646000	161.288000	
per_capita_gdp	mean	51189.000000	52681.250000	54014.250000	
	max	51986.000000	53242.000000	54317.000000	
	min	50462.000000	52179.000000	53719.000000	
urban_population	mean	79.583000	79.757000	79.928000	
	max	79.583000	79.757000	79.928000	
	min	79.583000	79.757000	79.928000	
year		2006	2007	2008	\
new_constructed_units	mean	1989.000000	1513.833333	1126.833333	
	max	2245.000000	1822.000000	1331.000000	
	min	1877.000000	1328.000000	1017.000000	
unemployment_rate	mean	4.608333	4.616667	5.800000	
	max	4.800000	5.000000	7.300000	
	min	4.400000	4.400000	4.900000	
interest_rates	mean	4.964167	5.019167	1.927500	
	max	5.250000	5.260000	3.940000	
	min	4.290000	4.240000	0.160000	
income	mean	11777.783333	12054.300000	12244.325000	
	max	11958.600000	12120.500000	12696.200000	
	min	11651.800000	11956.600000	12100.800000	
home_price_index	mean	183.447500	179.918917	164.057417	
	max	184.364000	184.598000	173.132000	
	min	182.321000	174.342000	153.619000	
per_capita_gdp	mean	54993.500000	55560.000000	55105.250000	
	max	55216.000000	55857.000000	55705.000000	
	min	54886.000000	55260.000000	53941.000000	
urban_population	mean	80.099000	80.269000	80.438000	
	max	80.099000	80.269000	80.438000	
	min	80.099000	80.269000	80.438000	
year		2009	2010	2011	\
new_constructed_units	mean	795.500000	653.750000	584.583333	
	max	846.000000	894.000000	634.000000	
	min	721.000000	552.000000	520.000000	
unemployment_rate	mean	9.283333	9.608333	8.933333	
	max	10.000000	9.900000	9.100000	
	min	7.800000	9.300000	8.500000	
interest_rates	mean	0.160000	0.175000	0.101667	
	max	0.220000	0.200000	0.170000	
	min	0.120000	0.110000	0.070000	
income	mean	12273.441667	12505.291667	12775.250000	
	max	12500.100000	12680.600000	12898.500000	
	min	12193.200000	12285.800000	12705.300000	
home_price_index	mean	148.545083	144.674500	139.259500	
	max	151.507000	147.396000	141.521000	
	min	147.694000	142.060000	136.674000	

per_capita_gdp	mean	53212.500000	54188.250000	54603.000000
	max	53531.000000	54569.000000	54979.000000
	min	53017.000000	53683.000000	54341.000000
urban_population	mean	80.606000	80.772000	80.944000
	max	80.606000	80.772000	80.944000
	min	80.606000	80.772000	80.944000
year		2012	2013	2014 \
new_constructed_units	mean	641.416667	763.166667	882.666667
	max	730.000000	839.000000	975.000000
	min	545.000000	708.000000	800.000000
unemployment_rate	mean	8.075000	7.358333	6.158333
	max	8.300000	8.000000	6.700000
	min	7.700000	6.700000	5.600000
interest_rates	mean	0.140000	0.107500	0.089167
	max	0.160000	0.150000	0.120000
	min	0.080000	0.080000	0.070000
income	mean	13125.383333	12937.191667	13383.675000
	max	13642.700000	13027.500000	13719.100000
	min	12961.600000	12813.900000	13079.500000
home_price_index	mean	140.993833	154.520750	164.698167
	max	145.503000	160.994000	168.050000
	min	136.533000	146.827000	161.927000
per_capita_gdp	mean	55422.250000	56171.750000	57137.750000
	max	55490.000000	56642.000000	57702.000000
	min	55342.000000	55859.000000	56345.000000
urban_population	mean	81.119000	81.299000	81.483000
	max	81.119000	81.299000	81.483000
	min	81.119000	81.299000	81.483000
year		2015	2016	2017 \
new_constructed_units	mean	965.250000	1060.500000	1151.833333
	max	1040.000000	1244.000000	1236.000000
	min	781.000000	945.000000	1069.000000
unemployment_rate	mean	5.275000	4.875000	4.358333
	max	5.700000	5.100000	4.700000
	min	5.000000	4.700000	4.100000
interest_rates	mean	0.132500	0.395000	1.001667
	max	0.240000	0.540000	1.300000
	min	0.110000	0.340000	0.650000
income	mean	13908.575000	14172.250000	14613.966667
	max	14060.200000	14313.700000	14797.900000
	min	13797.700000	14095.500000	14373.700000
home_price_index	mean	172.181750	180.925500	191.397667
	max	176.543000	185.722000	197.172000
	min	168.634000	177.274000	186.805000
per_capita_gdp	mean	58363.250000	58967.750000	60000.750000

	max	58486.00000	59296.0000	60674.000000	
	min	58121.00000	58704.0000	59494.000000	
urban_population	mean	81.67100	81.8620	82.058000	
	max	81.67100	81.8620	82.058000	
	min	81.67100	81.8620	82.058000	
year		2018	2019	2020	\
new_constructed_units	mean	1190.000000	1260.666667	1286.333333	
	max	1288.000000	1334.000000	1446.000000	
	min	1053.000000	1150.000000	1171.000000	
unemployment_rate	mean	3.891667	3.683333	8.091667	
	max	4.100000	4.000000	14.700000	
	min	3.700000	3.500000	3.500000	
interest_rates	mean	1.831667	2.158333	0.375833	
	max	2.270000	2.420000	1.580000	
	min	1.410000	1.550000	0.050000	
income	mean	15143.616667	15608.908333	16607.466667	
	max	15506.500000	15778.200000	18020.200000	
	min	14886.600000	15501.700000	15696.300000	
home_price_index	mean	202.476417	209.463333	222.143417	
	max	206.156000	213.933000	236.486000	
	min	198.315000	206.539000	214.994000	
per_capita_gdp	mean	61417.500000	62605.500000	60984.750000	
	max	61622.000000	63257.000000	62414.000000	
	min	61093.000000	61889.000000	57386.000000	
urban_population	mean	82.256000	82.459000	82.664000	
	max	82.256000	82.459000	82.664000	
	min	82.256000	82.459000	82.664000	
year		2021	2022	2023	
new_constructed_units	mean	1340.666667	1387.666667	1451.6000	
	max	1459.000000	1543.000000	1577.0000	
	min	1232.000000	1256.000000	1334.0000	
unemployment_rate	mean	5.366667	3.641667	3.6200	
	max	6.300000	4.000000	3.9000	
	min	3.900000	3.500000	3.4000	
interest_rates	mean	0.080000	1.683333	4.9630	
	max	0.100000	4.100000	5.3300	
	min	0.060000	0.080000	4.3300	
income	mean	17138.716667	16117.050000	16763.0200	
	max	20422.600000	16265.100000	16848.7000	
	min	16418.500000	15963.400000	16601.9000	
home_price_index	mean	260.045667	298.486750	303.8848	
	max	281.342000	304.724000	311.1750	
	min	239.560000	285.924000	297.0300	
per_capita_gdp	mean	64412.500000	65414.500000	66558.9000	
	max	65651.000000	65783.000000	67083.0000	

	min	63227.000000	65127.000000	66078.0000
urban_population	mean	82.873000	83.084000	83.0840
	max	82.873000	83.084000	83.0840
	min	82.873000	83.084000	83.0840

1. The number of new constructed units has generally increased over the years, with a noticeable peak in 2005 and a subsequent decline until around 2009.
2. The unemployment rate shows a pattern of decrease from 2003 to 2007, followed by a significant increase during the 2008 financial crisis.
3. Interest rates have experienced a general decline over the years, reaching a minimum around 2009.
4. The maximum income has consistently increased, reflecting overall economic growth.

1 Exploratory data analysis (EDA)

```
[21]: cols = ['new_constructed_units', 'unemployment_rate', 'interest_rates',
              'income', 'per_capita_gdp', 'urban_population']
```

```
[22]: fig = plt.figure(figsize = (60,60))

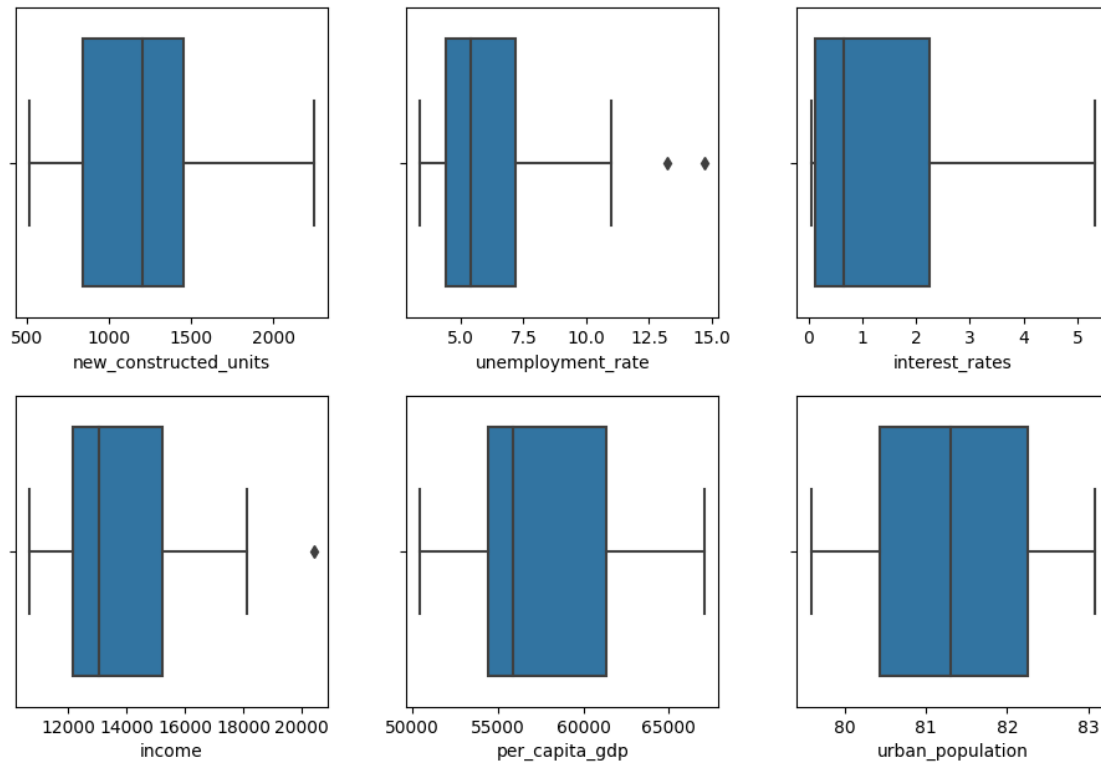
gs = fig.add_gridspec(15,15)

gs.update(wspace = 0.25,hspace = 0.25)

ax0 = fig.add_subplot(gs[0,0],)
ax1 = fig.add_subplot(gs[0,1])
ax2 = fig.add_subplot(gs[0,2])
ax3 = fig.add_subplot(gs[1,0])
ax4 = fig.add_subplot(gs[1,1])
ax5 = fig.add_subplot(gs[1,2])

Axis = [ax0,ax1,ax2,ax3,ax4,ax5]

for ax,col in zip(Axis,cols):
    sns.boxplot(ax = ax,data = df_20, x = col)
plt.show()
```




1. Income has 1 outliers greater than upper bound (19853.34) or lower than lower bound(7509.24). Cap them or remove them.
2. Unemployment Rate has 3 outliers greater than upper bound (11.40) or lower than lower bound(0.20). Cap them or remove them.

So below we are calculating the IQR and on the basis of lower and Upper limit we are capping the outliers.

```
[23]: sns.boxplot(data=df_20, x= 'income')
```

```
[23]: <Axes: xlabel='income'>
```



```
[24]: q1_income = df_20['income'].quantile(0.25)
      q3_income = df_20['income'].quantile(0.75)
      income_IQR = q3_income - q1_income

      upper_limit_income = q3_income + 1.5 * income_IQR
      lower_limit_income = q1_income - 1.5 * income_IQR

      print(upper_limit_income)
      print(lower_limit_income)
```

```
19853.337499999998
7509.2375
```

```
[25]: df_20['income'] = df_20['income'].clip(lower=lower_limit_income,
      ↪upper=upper_limit_income)
```


```
<ipython-input-25-2109b2de28d0>:1: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
df_20['income'] = df_20['income'].clip(lower=lower_limit_income,  
upper=upper_limit_income)
```

```
[26]: sns.boxplot(data=df_20, x= 'income')
```


```
[26]: <Axes: xlabel='income'>
```



We can see now there are no outliers present in the income data.

```
[27]: sns.boxplot(data=df_20, x= 'unemployment_rate')
```

```
[27]: <Axes: xlabel='unemployment_rate'>
```



```
[28]: q1_ur = df_20['unemployment_rate'].quantile(0.25)
      q3_ur = df_20['unemployment_rate'].quantile(0.75)
      ur_IQR = q3_ur - q1_ur

      upper_limit_ur = q3_ur + 1.5 * ur_IQR
      lower_limit_ur = q1_ur - 1.5 * ur_IQR

      print(upper_limit_ur)
      print(lower_limit_ur)
```

```
11.399999999999999
0.200000000000000107
```

```
[29]: df_20['unemployment_rate'] = df_20['unemployment_rate'].
      ↪clip(lower=lower_limit_ur, upper=upper_limit_ur)
```


```
<ipython-input-29-d67fd013cfd2>:1: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: <https://pandas.pydata.org/pandas->

```
docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
df_20['unemployment_rate'] =
df_20['unemployment_rate'].clip(lower=lower_limit_ur, upper=upper_limit_ur)
```


```
[30]: sns.boxplot(data=df_20, x= 'unemployment_rate')
```

```
[30]: <Axes: xlabel='unemployment_rate'>
```




We can see now there are no outliers present in the Unemployment Rate data.

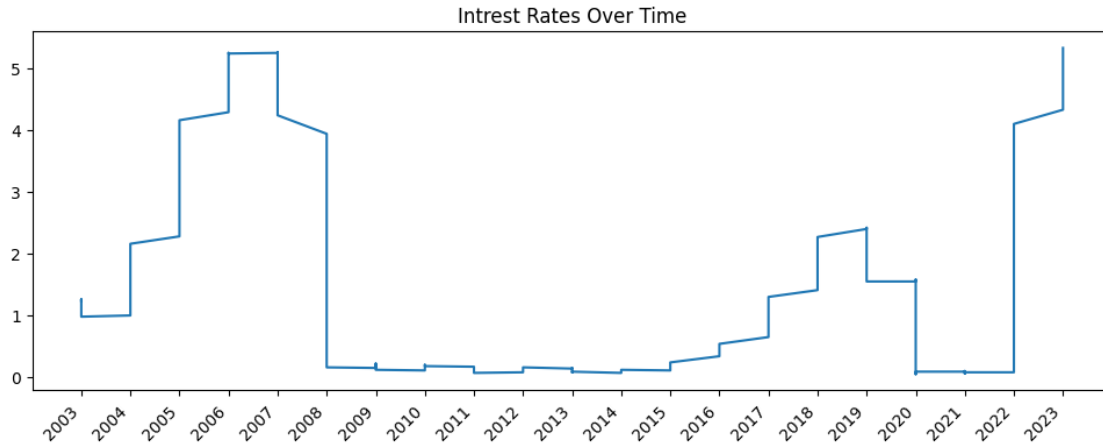
```
[31]: plt.figure(figsize = (12,4))
plt.plot(df_20['year'], df_20['new_constructed_units'])
plt.title("New Constructed Units Over Time")
plt.xticks(df_20['year'].unique(), rotation=45, ha="right")
plt.show()
```




```
[32]: plt.figure(figsize = (12,4))
plt.plot(df_20['year'], df_20['unemployment_rate'])
plt.title("Unemployment Rate Over Time")
plt.xticks(df_20['year'].unique(), rotation=45, ha="right")
plt.show()
```




```
[33]: plt.figure(figsize = (12,4))
plt.plot(df_20['year'], df_20['interest_rates'])
plt.title("Intrest Rates Over Time")
plt.xticks(df_20['year'].unique(), rotation=45, ha="right")
plt.show()
```




```
[34]: plt.figure(figsize = (12,4))
plt.plot(df_20['year'], df_20['income'])
plt.title("Income Over Time")
plt.xticks(df_20['year'].unique(), rotation=45, ha="right")
plt.show()
```




```
[35]: plt.figure(figsize = (12,4))
plt.plot(df_20['year'], df_20['per_capita_gdp'])
plt.title("Per Capita GDP Over Time")
plt.xticks(df_20['year'].unique(), rotation=45, ha="right")
plt.show()
```

```
[36]: plt.figure(figsize = (12,4))
plt.plot(df_20['year'], df_20['urban_population'])
plt.title("Urban Population Over Time")
plt.xticks(df_20['year'].unique(), rotation=45, ha="right")
plt.show()
```



```
[37]: plt.figure(figsize = (12,4))
plt.plot(df_20['year'], df_20['home_price_index'])
plt.title("Home Price Index Over Time")
plt.xticks(df_20['year'].unique(), rotation=45, ha="right")
plt.show()
```



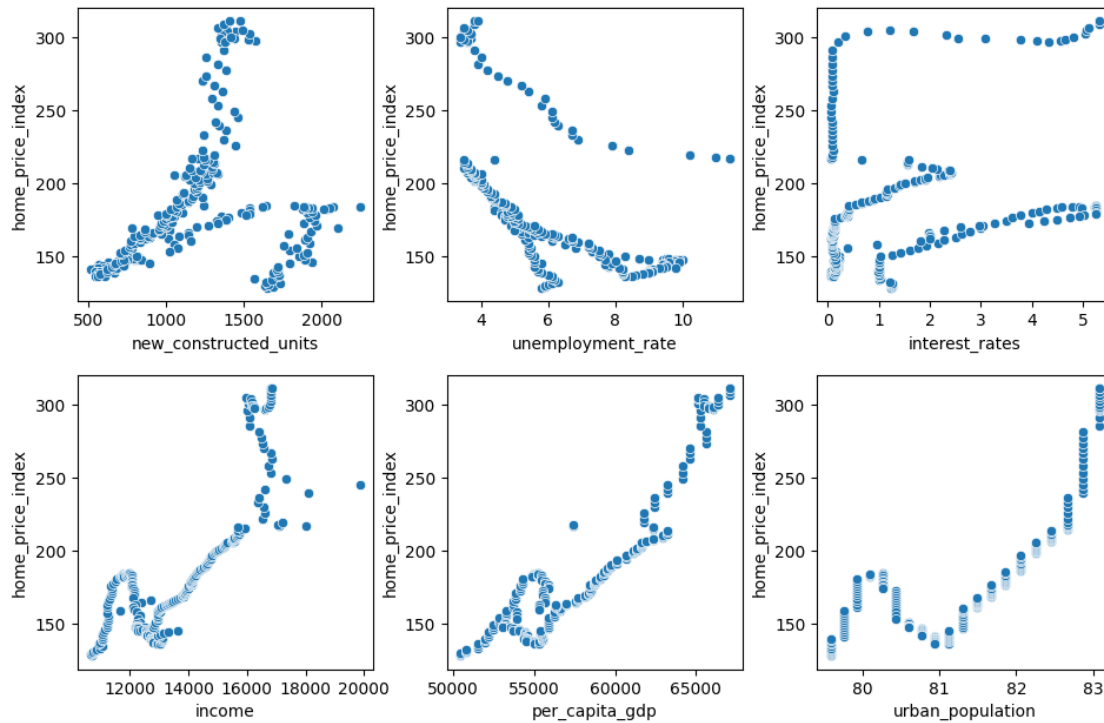
```
[38]: fig = plt.figure(figsize = (60,60))
gs = fig.add_gridspec(15,15)

gs.update(wspace = 0.25,hspace = 0.25)

ax0 = fig.add_subplot(gs[0,0],)
ax1 = fig.add_subplot(gs[0,1])
ax2 = fig.add_subplot(gs[0,2])
ax3 = fig.add_subplot(gs[1,0])
ax4 = fig.add_subplot(gs[1,1])
ax5 = fig.add_subplot(gs[1,2])

Axis = [ax0,ax1,ax2,ax3,ax4,ax5]

for ax,col in zip(Axis,cols):
    sns.scatterplot(df_20, ax=ax ,x = col, y = 'home_price_index')
plt.show()
```



“Scatter Plots of Factors vs. Home Price Index”

```
[39]: fig = plt.figure(figsize = (60,60))

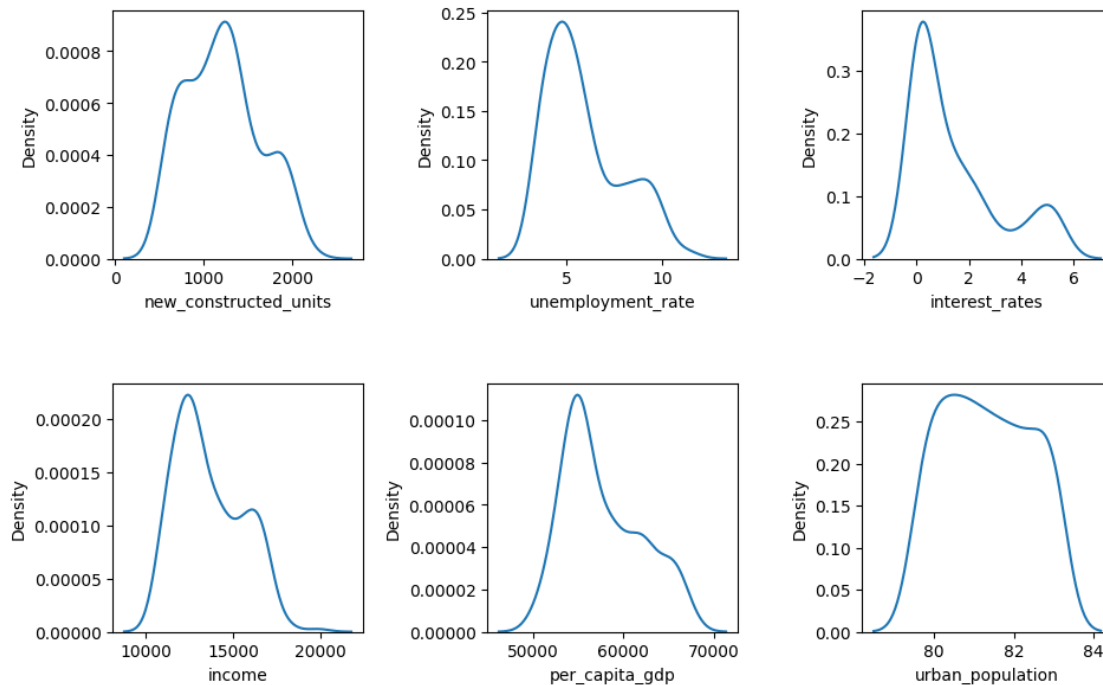
gs = fig.add_gridspec(15,15)

gs.update(wspace = 0.50,hspace = 0.50)

ax0 = fig.add_subplot(gs[0,0],)
ax1 = fig.add_subplot(gs[0,1])
ax2 = fig.add_subplot(gs[0,2])
ax3 = fig.add_subplot(gs[1,0])
ax4 = fig.add_subplot(gs[1,1])
ax5 = fig.add_subplot(gs[1,2])

Axis = [ax0,ax1,ax2,ax3,ax4,ax5]

for ax,col in zip(Axis,cols):
    sns.kdeplot(df_20, ax=ax ,x = col)
plt.show()
```



The variables `new_constructed_units` and `urban_population` exhibit Normal distributions, while the remaining variables show slight right-skewness. This mild skewness, although present, is deemed acceptable for model building, and we will proceed with our analysis considering these distribution characteristics as Normal.

```
[64]: df.corr()
```

```
<ipython-input-64-2f6f6606aa2c>:1: FutureWarning: The default value of
numeric_only in DataFrame.corr is deprecated. In a future version, it will
default to False. Select only valid columns or specify the value of numeric_only
to silence this warning.
```

```
df.corr()
```

```
[64]:
```

	new_constructed_units	unemployment_rate	\
new_constructed_units	1.000000	-0.390979	
unemployment_rate	-0.390979	1.000000	
interest_rates	0.418073	0.066943	
income	-0.411212	-0.046088	
home_price_index	-0.062159	-0.231612	
per_capita_gdp	-0.376994	0.103897	
urban_population	-0.418469	-0.068268	

	interest_rates	income	home_price_index	\
new_constructed_units	0.418073	-0.411212	-0.062159	
unemployment_rate	0.066943	-0.046088	-0.231612	

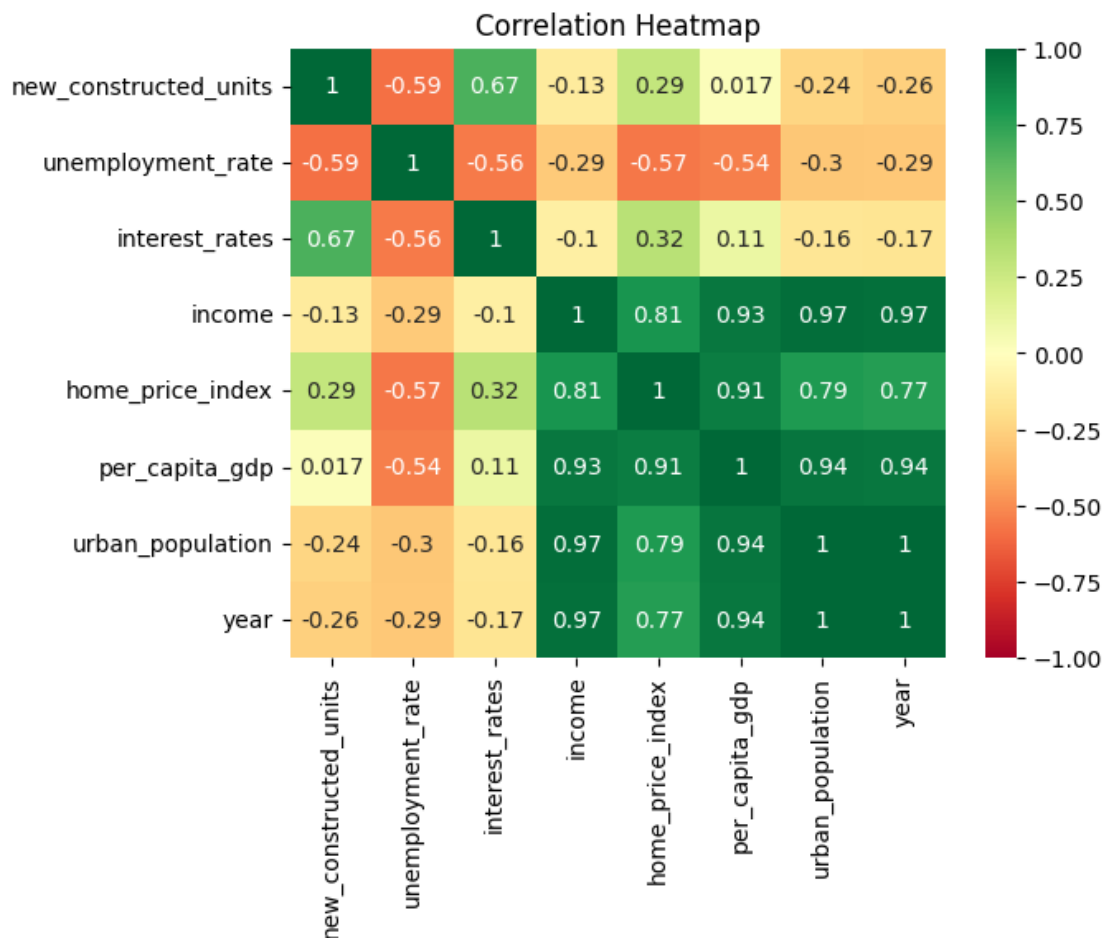
interest_rates	1.000000	-0.541979	-0.542286
income	-0.541979	1.000000	0.935364
home_price_index	-0.542286	0.935364	1.000000
per_capita_gdp	-0.374240	0.987732	0.943785
urban_population	-0.552013	0.986303	0.889301

	per_capita_gdp	urban_population
new_constructed_units	-0.376994	-0.418469
unemployment_rate	0.103897	-0.068268
interest_rates	-0.374240	-0.552013
income	0.987732	0.986303
home_price_index	0.943785	0.889301
per_capita_gdp	1.000000	0.990761
urban_population	0.990761	1.000000

```
[40]: sns.heatmap(df_20.corr(),vmax = 1, vmin = -1 , annot = True , cmap = "RdYlGn")
plt.title("Correlation Heatmap")
plt.show()
```

<ipython-input-40-748b961f0c4e>:1: FutureWarning: The default value of numeric_only in DataFrame.corr is deprecated. In a future version, it will default to False. Select only valid columns or specify the value of numeric_only to silence this warning.

```
sns.heatmap(df_20.corr(),vmax = 1, vmin = -1 , annot = True , cmap = "RdYlGn")
```



The above Heatmap indicates that few of these variables provide similar information, and their high correlation might lead to multicollinearity issues in regression analysis.

2 Standardization

Standardization is performed on the dataset to bring all features to a common scale. This is crucial when working with machine learning models that are sensitive to the magnitude of input variables.

```
[41]: from sklearn.model_selection import train_test_split
      from sklearn.preprocessing import StandardScaler
      scaler = StandardScaler()
```

```
[42]: X = df_20.drop(columns = ['home_price_index', 'date', 'year'])
      y = df_20["home_price_index"]
```

```
[43]: X_train,X_test,y_train,y_test = train_test_split(X,y,test_size= 0.33,
      ↪random_state=2)
```

```
[44]: X_train_scaled = scaler.fit_transform(X_train)
      X_test_scaled = scaler.transform(X_test)
```

3 MODEL BUILDING

4 Linear Regression

```
[45]: from sklearn.linear_model import LinearRegression
      lr = LinearRegression()
```

```
[46]: lr.fit(X_train_scaled,y_train)
      lr_pred = lr.predict(X_test_scaled)
```

```
[47]: r2_lr = r2_score(y_test,lr_pred)
      print("R squared Value")
      print(r2_score(y_test,lr_pred))
      mse_lr = mean_squared_error(y_test,lr_pred)
      print("\nMSE")
      print(mean_squared_error(y_test,lr_pred))
      mae_lr = mean_absolute_error(y_test,lr_pred)
      print("\nMAE")
      print(mean_absolute_error(y_test,lr_pred))
```

R squared Value
0.9527354255690559

MSE
119.27608462812941

MAE
7.788883975430829

R2 Score for LR model indicates 95.27% of the variability in the dependent variable is explained by the model.

5 Random forest Model

```
[48]: from sklearn.ensemble import RandomForestRegressor
      rr = RandomForestRegressor()
```

```
[49]: rr.fit(X_train_scaled,y_train)
      rr_pred = rr.predict(X_test_scaled)
```

```
[50]: r2_rr = r2_score(y_test,rr_pred)
      print("R squared Value")
      print( r2_score(y_test,rr_pred))
```

```

mse_rr = mean_squared_error(y_test,rr_pred)
print("\nMSE")
print( mean_squared_error(y_test,rr_pred))
mae_rr = mean_absolute_error(y_test,rr_pred)
print("\nMAE")
print( mean_absolute_error(y_test,rr_pred))

```

R squared Value
0.9962608494326256

MSE
9.43605744643026

MAE
2.0383102409638623

R2 Score for RR model indicates 99.65% of the variability in the dependent variable is explained by the model.

6 ADA BOOST

```

[51]: from sklearn.ensemble import AdaBoostRegressor
ada = AdaBoostRegressor()

```

```

[52]: ada.fit(X_train_scaled,y_train)
ada_pred = ada.predict(X_test_scaled)

```

```

[53]: r2_ab = r2_score(y_test,ada_pred )
print("R squared Value")
print( r2_score(y_test,ada_pred ))
mse_ab = mean_squared_error(y_test,ada_pred )
print("\nMSE")
print( mean_squared_error(y_test,ada_pred ))
mae_ab = mean_absolute_error(y_test,ada_pred )
print("\nMAE")
print( mean_absolute_error(y_test,ada_pred ))

```

R squared Value
0.9900592689000133

MSE
25.086261713407538

MAE
3.834583774110322

R2 Score for ADA Boosting model indicates 98.84% of the variability in the dependent variable is explained by the model.

7 Gradient Boosting

```
[54]: from sklearn.ensemble import GradientBoostingRegressor
gb = GradientBoostingRegressor()
```

```
[55]: gb.fit(X_train_scaled, y_train)
gb_pred = gb.predict(X_test_scaled)
```

```
[56]: r2_gb = r2_score(y_test,gb_pred )
print("R squared Value")
print( r2_score(y_test,gb_pred ))
mse_gb = mean_squared_error(y_test,gb_pred )
print("\nMSE")
print( mean_squared_error(y_test,gb_pred ))
mae_gb = mean_absolute_error(y_test,gb_pred )
print("\nMAE")
print( mean_absolute_error(y_test,gb_pred ))
```

R squared Value
0.9947513910680241

MSE
13.245301172974209

MAE
2.6471187223910935

R2 Score for Gradient Boosting model indicates 99.48% of the variability in the dependent variable is explained by the model.

8 XG Boosting Model

```
[57]: from xgboost import XGBRegressor
xg = XGBRegressor()
```

```
[58]: xg.fit(X_train_scaled, y_train)
xg_pred = xg.predict(X_test_scaled)
```

```
[59]: r2_xg = r2_score(y_test,xg_pred )
print("R squared Value")
print( r2_score(y_test,xg_pred ))
mse_xg = mean_squared_error(y_test,xg_pred )
print("\nMSE")
print( mean_squared_error(y_test,xg_pred ))
mae_xg = mean_absolute_error(y_test,xg_pred )
print("\nMAE")
print( mean_absolute_error(y_test,xg_pred ))
```

R squared Value
0.9953096425659546

MSE
11.836507087496674

MAE
2.3933190462043488

R2 Score for ADA Boosting model indicates 99.53% of the variability in the dependent variable is explained by the model.


9 Observation

Random Forest and XGBoosting consistently outperform other models, demonstrating their effectiveness in capturing the underlying patterns in the data. Below is the graph that shoes the actual and the predicted value for the Ada Boost Algorithmn.

```
[63]: fit_data = ada.fit(X_train_scaled,y_train)
      prediction = fit_data.predict(X_test_scaled)

      plt.figure(figsize=(10,10))
      plt.scatter(y_test, prediction, c = '#d62728')
      plt.yscale('log')
      plt.xscale('log')

      p1 = max(max(prediction), max(y_test))
      p2 = min(min(prediction), min(y_test))
      plt.plot([p1, p2], [p1, p2], 'b-')
      plt.xlabel('True Values', fontsize=15)
      plt.ylabel('Predictions', fontsize=15)
      plt.title("Accuracy Plot")
      plt.axis('equal')
      plt.show()
```



[]: